

DSAPS Assignment 2

Deadline: 11 October 2026, 11:59pm

Important Points

1. The assignment has 3 questions. You need to do ALL THREE questions. Each question weighs 100 marks.
2. Only C++ is allowed.
3. Directory structure to be followed for submission:

```
202620xxxx_A2
|--202620xxxx_A2_Q1A.cpp
|--202620xxxx_A2_Q1B.cpp
|--202620xxxx_A2_Q2A.cpp
|--202620xxxx_A2_Q2B.cpp
|--202620xxxx_A2_Q3.cpp
|--README.md
```

Replace your roll number in place of 202620xxxx.

4. **Submission Format:** Follow the above mentioned directory structure and zip the RollNo_A2 folder and submit RollNo_A2.zip on Moodle.

Note: All submissions which are not in the specified format will be awarded 0 in the assignment.

5. C++ STL (including vectors) is not allowed for any of the questions unless specified otherwise in the question. So `#include <bits/stdc++.h>` is not allowed.
6. A brief description of your approach for solving each question (data structures, algorithms, and any optimizations used) and step-by-step instructions on how to compile and execute your code for each question should be mentioned in README.
7. You can ask queries by posting on Moodle.
8. **Late Submission Rule:** Deadlines for assignments are final and will not be extended. Late submissions cost 5% loss of marks for each late day up to 3 days (i.e., 5%, 10%, ...). Submissions beyond 3 days from the deadline shall receive 0 marks.

NOTE: In case of plagiarism in any of the questions, the students involved will be awarded -10 as the final marks of Assignment 2.

1 Question 1. Deque

1.1 A. [80 marks]

Problem Statement: Implement Deque

What is deque?

- Deque is the same as dynamic arrays with the ability to resize itself automatically when an element is inserted, with their storage being handled automatically by the container.
- They support insertion and deletion from both ends in amortized constant time.
- Inserting and erasing in the middle is linear in time.

Operations: The C++ standard specifies that a legal (i.e., standard-conforming) implementation of deque must satisfy the following performance requirements: (consider the data type as T)

1. `deque()` - initialize an empty deque. Time complexity: $O(1)$
2. `deque(n)` - initialize a deque of length n with all values as default value of T . Time complexity: $O(n)$
3. `deque(n, x)` - Initialize a deque of length n with all values as x . Time complexity: $O(n)$
4. `bool push_back(x)` - append data x at the end. Return true if operation is performed successfully, else return false. Time complexity: constant amortized time
5. `bool pop_back()` - erase data at the end. Return true if operation is performed successfully, else return false. Time complexity: constant amortized time
6. `bool push_front(x)` - append data x at the beginning. Return true if operation is performed successfully, else return false. Time complexity: constant amortized time
7. `bool pop_front()` - erase an element from the beginning. Return true if operation is performed successfully, else return false. Time complexity: constant amortized time
8. `T front()` - returns the first element(value) in the deque. If the first element is not present, return the default value of T . Time complexity: $O(1)$
9. `T back()` - returns the last element(value) in the deque. If the last element is not present, return the default value. Time complexity: $O(1)$
10. `T D[n]` - returns the n th element of the deque. You need to overload the `[]` operator. If n th element is not present return default value of T . Time complexity: $O(1)$
11. `bool empty()` - returns true if deque is empty else returns false. Time complexity: $O(1)$
12. `int size()` - returns the current size of deque. Time complexity: $O(1)$
13. `void resize(n)` - change the size dynamically to new size n . Time complexity: $O(n)$
 - If the new size n is greater than the current size of the deque, then insert new elements with the default value of T at the end of the queue.
 - If the new size n is smaller than the current size, then keep n elements from the beginning of the deque.
14. `void resize(n, d)` - change the size dynamically to new size n . Time complexity: $O(n)$

- If the new size n is greater than the current size of the deque, then insert new elements with value d at the end of the queue.
 - If the new size n is smaller than the current size, then keep n elements from the beginning of the deque.
15. `void reserve(n)` : change the capacity of deque to n , if $n >$ current capacity; otherwise do nothing. Time complexity: $O(n)$
 16. `void shrink_to_fit()` - reduce the capacity of the deque to current size. Time Complexity: $O(\text{size}())$
 17. `void clear()` - remove all elements of deque. Time complexity: $O(n)$
 18. `int capacity()` - return the current capacity of deque. Time complexity: $O(1)$

Note:

1. Your deque should be generic type i.e. it should be datatype independent and can support primitive data types like integer, float, string, etc. Hint: Use template in C++ Hint: Use template in C++ (Link).
2. For operations 1, 2 & 3 you can either define a constructor for the class or initialize the class object using void return type functions.
3. C++ STL is not allowed (including vectors, design your own if required)
4. `D[0]` - element at index 0 (i.e. first element from the front), `D[1]` - element at index 1 (i.e. second element from the front), `D[-1]` - element at last index (i.e. first element from the back), `D[-2]` - element at second last index (i.e. second element from the back)
5. Size of the deque is the number of elements currently present in your deque.
6. Capacity of the deque is the number of elements your deque can accommodate with currently held memory.
7. During Operation 1 both size and capacity of the deque should be set to zero.
8. If size is equal to capacity and a new element is inserted, then the capacity is doubled, unless capacity is zero, then it will become one.
9. If you have doubts about deciding the new capacity in any of the operations, refer to the behavior of the member functions of STL vector containers.

Input Format: Design an infinitely running menu-driven `main` function. Each time the user inputs an integer corresponding to the serial number of the operation listed above in the assignment. Then, take necessary arguments related to the selected operation and execute the respective method. Finally, the program must exit with status code 0, when 0 is provided as a choice.

Output Format: Each output should be in a new line. Use `boolalpha` to output booleans, e.g.:
`cout << boolalpha << dq.pop_front() << endl;`

Sample input:

```
1
4 10
```

```
6 20
12
8
5
11
0
```

Sample output:

```
true
true
2
20
true
false
```

1.2 B. [20 marks]

Problem Statement: Using the deque implemented above, you are supposed to implement a randomized queue. Implement the following operations:

1. `void enqueue(x)`: add the item to the randomized queue
2. `T dequeue()`: remove and return a random item
3. `T sample()`: return a random sample (but do not remove it)

The above operations must be done in constant amortized time. Randomness of fetched items will be checked.

Randomized queue: A randomized queue is similar to a stack or queue, except that the item removed is chosen uniformly at random from items in the data structure.

Constant amortized time: Refer to this Hint: Use template in C++ ([Link](#)).

Note: Each element in the randomized queue must have an equal probability of being selected. Only these three operations will be checked in this part of the question.

Input Format: Design an infinitely running menu-driven `main` function. Each time the user inputs an integer corresponding to the serial number of the operation listed above. Then, take necessary arguments related to the selected operation and execute the respective method. Finally, the program must exit with status code 0, when 0 is provided as a choice.

2 Question 2. Priority Queue

2.1 A. [80 marks]

Problem Statement: Implement priority queue

What is a priority queue?

- A priority queue is a data structure that maintains a collection of elements, each associated with a priority or value. Elements are stored in a way that allows the retrieval of the element with the highest (or lowest) priority quickly.
- Unlike dequeues or lists, a priority queue doesn't maintain the elements in any specific order based on their values, except for ensuring that the highest (or lowest) priority element is readily accessible.

Operations: The C++ standard specifies that a legal (i.e., standard-conforming) implementation of priority queue must satisfy the following performance requirements.

1. `priority_queue()` - initialize an empty priority queue. Time complexity: $O(1)$
2. `int size()` - returns the current size of the priority queue. Time complexity: $O(1)$
3. `void push(int el)` - insert an element `el` in the priority queue. Time complexity: $O(\log(\text{size}()))$
4. `int top()` - returns the top (highest or lowest priority) element in the priority queue. Time complexity: $O(1)$
5. `void pop()` - remove the top element of the priority queue. Time complexity: $O(\log(\text{size}()))$
6. `bool empty()` - returns true if the priority queue is empty else returns false. Time complexity: $O(1)$

2.2 B. [20 marks]

Problem Statement:

David's Bakery has introduced a new promotion for students. Under this promotion, he monitors his daily sales, and if the sales on a specific day is greater than or equal to the combined sum of the median sale for a trailing number of d days and the median sale from the first day of the promotion, (follow through the example included below to understand better) then he offers free 'Cheeeese Maggi' to the first five customers on the following day. However, David will only provide this free offer if he has transaction data for at least the trailing number of d prior days. Given the number of trailing days d and the daily sales record for David's Bakery over a period of n days, can you determine how many days David will offer free maggi to students during this entire period of n days?

Input Format:

- The first line contains two space-separated integers n and d , the number of days of daily sales data, and the number of trailing days respectively.
- The second line contains n space-separated non-negative integers where each integer i denotes `sales[i]`.

Output Format:

- **int**: the total number of days David will offer free maggi to students.

Constraints:

$$1 \leq n \leq 2 \times 10^5$$

$$1 \leq d < n$$

$$1 \leq \text{sales}[i] \leq 10^5$$

Sample input:

```
8 5
2 3 4 2 1 6 3 6
```

Sample output:

```
2
```

Explanation:

For the first five days, David will not offer any free maggi to students because he didn't have sufficient sales data.

On the sixth day, David has $d = 5$ days worth of trailing sales data, which is 2, 3, 4, 2, 1. The median sale of the preceding five days is 2, and he also has the entire sales data from the beginning, which remains 2, 3, 4, 2, 1, with a median sale value of 2 as well. Since David's sales on the 6th day is greater than the combined sum, David offers free Maggi to students for the 6th day.

On the seventh day, David has $d = 5$ days worth of trailing sales data, which is 3, 4, 2, 1, 6. The median sale of the preceding five days is 3, and he also has the entire sales data from the beginning, which is 2, 3, 4, 2, 1, 6, with a median sale value of $(2 + 3)/2 = 2.5$. Since David's sales on the 7th day is less than the combined sum, he will not offer free Maggi to students for the 7th day.

On the eighth day, David has $d = 5$ days worth of trailing sales data, which is 4, 2, 1, 6, 3. The median sale of the preceding five days is 3, and he also has the entire sales data from the beginning, which is 2, 3, 4, 2, 1, 6, 3, with a median sale value of 3 as well. Since David's sales on the 8th day is equal to the combined sum, David offers free Maggi to students for the 8th day.

Therefore, in total David will offer free maggi for 2 days (6th and 8th day).

Note: You need to implement priority queue for Integer data type only. You need to implement both minimum and maximum priority queue (Implementation details are up to you).


3 Question 3: The Lost Library of IIIT Hyderabad

Problem

There is a row of n shelves (“bins”), numbered 1 to n left to right, each holding one unique book. Because the books are sorted by ID, you can always find a book with a binary search: pick some bin to check first, and depending on whether the target is smaller or larger, go left or right, and so on.

Hint: A binary search over n sorted bins is a **binary search tree (BST)** on the bins $1, \dots, n$: whichever bin you check first is the root, everything to its left forms the left subtree, and everything to its right forms the right subtree. The number of bins you check to find bin i is simply the *depth* of bin i in this tree (root = depth 1).

Each book i is looked up $F[i]$ times a month, so books that are asked for often should sit near the root (shallow), and rarely-asked books can sit deeper. Your job in all three parts is to choose the shape of this search tree so that some notion of “total cost” is as small as possible. The three parts just change what “cost” means.



3.1 Part 1: The Ideal Search Plan [30 marks]

Ignore movement for now: the only cost is *how many bins get checked*. If book i sits at depth $\text{depth}(i)$ in your search tree, finding it costs $\text{depth}(i)$ checks, and this happens $F[i]$ times a month. So the total monthly cost of a tree is

$$\text{Cost} = \sum_{i=1}^n F[i] \cdot \text{depth}(i).$$

Your task: choose the shape of the BST (i.e., decide which bin to check first, then which bin to check next in each half, and so on) so that this total cost is as small as possible.

Input: An array $F[1 \dots n]$, where $F[i]$ is the monthly request frequency for the book in bin i .

Worked example: Let $n = 3$, $F = [10, 2, 5]$ (see the shelves above). The naive strategy always checks the middle bin first - bin 2, then bin 1 or bin 3 (the tree on the right above):

- Book 2 sits at depth 1: cost $= 1 \times F[2] = 1 \times 2 = 2$.
- Book 1 sits at depth 2: cost $= 2 \times F[1] = 2 \times 10 = 20$.
- Book 3 sits at depth 2: cost $= 2 \times F[3] = 2 \times 5 = 10$.

Total cost of this naive plan: $2 + 20 + 10 = 32$. Since book 1 is asked for the most (frequency 10) but sits deep in this tree, this plan wastes time. A better tree would keep book 1 shallow instead. This shows that always picking the middle bin is not always optimal. Your task is to find a tree that beats this naive approach and finds the true minimum possible total cost.

3.2 Part 2: Introducing Basic Robot Costs [30 marks]


Now the robot is real: it physically drives along the shelf, starting from a charging station at position 0, and every bin it visits along the way costs distance, not just “one check”. So instead of counting checks, we count how far the robot actually travels.

New Cost Model:

- Moving from bin i to bin j costs $|i - j|$ seconds (just the distance).
- The search ends the moment the target bin is scanned — no need to go anywhere else yet.

So the cost of finding book i is now the total distance the robot drives while following the same left/right search path as before (root, then its child, etc.), not just the number of bins on that path. The overall goal is unchanged in spirit: minimize $\sum_i F[i] \cdot (\text{distance traveled to reach bin } i)$, but now the arrangement of the tree also determines physical distances, since bins that are far apart on the shelf cost more to hop between.

Worked example: Using the same naive plan (check bin 2, then bin 1 or bin 3):



- Book 2: drive $0 \rightarrow 2$, distance 2. Contribution $= 2 \times F[2] = 4$.
- Book 1: drive $0 \rightarrow 2 \rightarrow 1$, distance $|2 - 0| + |1 - 2| = 3$. Contribution $= 3 \times F[1] = 30$.
- Book 3: drive $0 \rightarrow 2 \rightarrow 3$, distance $|2 - 0| + |3 - 2| = 3$. Contribution $= 3 \times F[3] = 15$.

Total cost of this plan: $4 + 30 + 15 = 49$.

Your Goal: Find the search plan that minimizes the total expected time with this new cost model.

Your Tasks for Part 2:

1. Adapt your approach from Part 1 to work with this more complex cost model. How does your recursive definition for the optimal cost change?
2. Describe an efficient algorithm to compute the minimum possible search time.
3. Analyze the time and space complexity of your algorithm.


3.3 Part 3: The Full Realistic Model [40 marks]

This part just adds three more real-world details on top of Part 2: moving costs a rate α per unit distance (not exactly 1 anymore), scanning a bin takes a fixed time β , changing direction mid-search wastes an extra γ seconds, and once the book is found the robot has to drive all the way back to the charging station.

Real Time Model:

- Moving from bin i to bin j costs $\alpha \cdot |i - j|$ seconds.
- Scanning the book ID at any bin costs a flat β seconds.
- Flipping direction (right-then-left or left-then-right) costs an extra γ seconds. E.g., driving $0 \rightarrow 5 \rightarrow 2$ flips direction once.
- After the book is found at bin i , the robot must drive back home, adding $\alpha \cdot i$ seconds (this final trip is never penalized for reversing).

Worked example: Same tree as before (check bin 2 first), with $\alpha = 2$, $\beta = 10$, $\gamma = 4$.



- **Book 2** (found immediately, no reversal): drive $0 \rightarrow 2$, scan, drive back $2 \rightarrow 0$: $(2 \cdot 2) + 10 + (2 \cdot 2) = 18$. Contribution = $18 \times F[2] = 36$.
- **Book 1** (one reversal, since the robot overshoots to bin 2 first): drive $0 \rightarrow 2$, scan, *reverse*, drive $2 \rightarrow 1$, scan, drive back $1 \rightarrow 0$: $(2 \cdot 2) + 10 + 4 + (2 \cdot 1) + 10 + (2 \cdot 1) = 32$. Contribution = $32 \times F[1] = 320$.
- **Book 3** (no reversal, robot keeps moving the same direction): drive $0 \rightarrow 2$, scan, drive $2 \rightarrow 3$, scan, drive back $3 \rightarrow 0$: $(2 \cdot 2) + 10 + (2 \cdot 1) + 10 + (2 \cdot 3) = 32$. Contribution = $32 \times F[3] = 160$.

Total cost of this plan: $36 + 320 + 160 = 516$.

Your Tasks for Part 3:

1. The reversal penalty γ is tricky because the cost of a move now depends on the previous move. Explain how you would need to change your recursive definition to handle this. What new information must your subproblem definition keep track of?
2. Describe the final, efficient algorithm that solves this complete problem.
3. Analyze the time and space complexity of this final algorithm.

Constraints

For all parts of the problem, you can assume the inputs will follow these constraints:

$$\begin{aligned}1 &\leq n \leq 80 \\1 &\leq F[i] \leq 1,000,000 \\1 &\leq \alpha, \beta, \gamma \leq 1,000\end{aligned}$$

- The total cost may exceed the capacity of a 32-bit integer.

Testing Your Algorithm: Input/Output Format

Your final program should be able to read the query type and corresponding parameters from standard input to solve the appropriate version of the problem.

Query Type 1 (for Part 1) Input Format:

```
1
n
F_1 F_2 ... F_n
```

Example Input:

```
1
3
10 2 5
```

Example Output:

```
26
```

Hint: Path: $1 \rightarrow 3 \rightarrow 2$

Query Type 2 (for Part 2) Input Format:

```
2
n
F_1 F_2 ... F_n
```

Example Input:

```
2
3
10 2 5
```

Example Output:

```
29
```

Hint: Path: $1 \rightarrow 2 \rightarrow 3$

Query Type 3 (for Part 3) Input Format:

```
3
n
F_1 F_2 ... F_n
alpha beta gamma
```

Example Input:

```
3
3
10 2 5
2 10 4
```

Example Output:

```
392
```

Hint: Path: $1 \rightarrow 3 \rightarrow 2$

Notes:

- The reversal penalty only applies to movements during the search phase, not for the final “return home” trip.
- You are not allowed to use `vector` and all, so if you require any function make your own.